

Math Bootcamp

Day 3 — Remaining real analysis and linear algebra

Math Camp 2026

How to use these slides

Main deck: definitions, recipes, examples, and practice.

Appendix and proofs (marked on relevant slides): **for interest only** — optional recap for those already comfortable with the math.

Not required for Math Camp. Focus on the main slides and exercises unless you want the extra detail.

3.1 Fixed points — equilibrium as a solution

Fixed point. A point x^* is a **fixed point** of a map T if $T(x^*) = x^*$.

Economics reading. Many equilibria are fixed-point problems:

- **Market clearing:** a price p^* with excess demand $E(p^*) = 0$.
- **Game theory:** a strategy profile that is a best response to itself (Nash).

Two useful existence tools (today).

- **Contraction Mapping Theorem:** existence and uniqueness of a fixed point; iteration converges.
- **Brouwer:** continuous map on a closed interval / convex set.

3.1 Contraction Mapping Theorem

Contraction. A map T **shrinks distances** if there is $k \in (0, 1)$ such that $|T(x) - T(y)| \leq k|x - y|$ for all x, y in the domain.

Contraction Mapping Theorem. If $T: [a, b] \rightarrow [a, b]$ is a contraction, then

- **Existence:** there is a fixed point x^* with $T(x^*) = x^*$.
- **Uniqueness:** that fixed point is **unique**.
- **Iteration:** for any starting value $x_0 \in [a, b]$, the sequence $x_n = T(x_{n-1})$ converges to x^* .

3.1 Why iteration converges

Step 1 — Start anywhere. Pick any x_0 . Define $x_n = T(x_{n-1})$ for $n = 1, 2, \dots$

Step 2 — Jumps shrink exponentially. Each new step is at most k times the previous jump:

$$|x_{n+1} - x_n| = |T(x_n) - T(x_{n-1})| \leq k|x_n - x_{n-1}| \leq k^n|x_1 - x_0|.$$

Since $k < 1$, the moves get **exponentially smaller** — the path barely changes after a while.

Step 3 — The sequence converges. Step 2 implies $x_n \rightarrow x^*$ for some limit x^* (check it is a **Cauchy sequence** for a rigorous proof). So iteration from **any** starting point drives x_n toward the same x^* .

Step 4 — The limit is a fixed point. T is continuous, so taking limits in $x_n = T(x_{n-1})$ gives $x^* = \lim x_n = \lim T(x_{n-1}) = T(\lim x_{n-1}) = T(x^*)$.

3.1 Walrasian price adjustment

Question. Excess demand $E(p)$, market clearing $E(p^*) = 0$. Prices update by

$$p_{t+1} = T(p_t) = p_t + \lambda(1 - e^{-\alpha E(p_t)}), \quad \lambda = \frac{1}{5}, \alpha = 0.1.$$

Does the price stabilize? That is, does $p_t \rightarrow p^*$?

Fixed points of T are exactly the market-clearing prices: $E(p^*) = 0 \Leftrightarrow T(p^*) = p^*$. Near $E = 0$, $1 - e^{-\alpha E} \approx \alpha E$.

Answer. On a compact range that T maps into itself, if $E' < 0$ and $T'(p) > -1$ (equivalently $\lambda \alpha e^{-\alpha E(p)} |E'(p)| < 2$), then yes: T is a contraction, so $p_t \rightarrow p^*$ from any start in the range.

3.1 Why the price stabilizes

Slope. $T'(p) = 1 + \lambda \alpha e^{-\alpha E(p)} E'(p)$.

Why $T' < 1$. $E' < 0$, so the second term is negative and $T' < 1$.


Why $T' > -1$ is enough. Then $-1 < T' < 1$, so $\max |T'| = k < 1$. By MVT,
 $|T(p) - T(q)| = |T'(c)| |p - q| \leq k |p - q|$: each update shrinks the gap to p^* , hence the price stabilizes.

3.1 Brouwer — one dimension

Brouwer (one dimension). If $T: [a, b] \rightarrow [a, b]$ is **continuous**, then T has **at least one** fixed point x^* with $T(x^*) = x^*$.

Picture. Plot $y = T(x)$ against the 45-degree line $y = x$. A crossing is a fixed point.

Multiple fixed points possible



3.1 Brouwer — general statement


Brouwer (several variables). Let $K \subset \mathbb{R}^n$ be **nonempty**, **compact**, and **convex**. If $T: K \rightarrow K$ is **continuous**, then T has at least one fixed point $x^* \in K$ with $T(x^*) = x^*$.

Camp vocabulary.

- **Compact** in $\mathbb{R}^n$: closed and bounded (Heine–Borel).
- **Convex**: the segment between any two points in K stays in K .

One dimension. $K = [a, b]$ is compact and convex, so the 1D case is a special case.

3.1 Why compact and convex



Left: $K = (0, 1)$ not closed, $T(x) = x/2$; the crossing is at $0 \notin K$. Right: K is two intervals, T swaps them; no crossing on K .

3.1 Linear Cournot — best response

Example. Two firms, $p = a - (q_1 + q_2)$, cost 0, $q_i \in [0, a]$, $a > 0$. Firm 1 takes q_2 as given: $\pi_1(q_1; q_2) = (a - q_1 - q_2)q_1$. Firm 2 is symmetric.

FOC. $\partial\pi_1/\partial q_1 = a - 2q_1 - q_2 = 0$, so $q_1 = (a - q_2)/2$. **SOC.** $\partial^2\pi_1/\partial q_1^2 = -2 < 0$.

Best response. $BR_1(q_2) = (a - q_2)/2$, $BR_2(q_1) = (a - q_1)/2$, and both map $[0, a]$ into itself.

3.1 Linear Cournot — Nash

Nash. $q_1^* = BR_1(q_2^*)$ and $q_2^* = BR_2(q_1^*)$:

$$2q_1^* = a - q_2^*, \quad 2q_2^* = a - q_1^*.$$

Substitute $q_2^* = a - 2q_1^*$ into the second equation: $2(a - 2q_1^*) = a - q_1^*$, so $2a - 4q_1^* = a - q_1^*$, hence $a = 3q_1^*$.

Thus $q_1^* = q_2^* = a/3$. Unique solution of two linear equations — no Brouwer needed.

Same map: $T(q_1, q_2) = (BR_1(q_2), BR_2(q_1))$ on $K = [0, a]^2$, with unique fixed point $(a/3, a/3)$.

3.1 Brouwer — Cournot duopoly

Asymmetric nonlinear demand. Two firms, $q_i \in [0, \bar{q}]$, $p = a - (q_1 + 2q_2)^2$ (e.g. $a = 8$, $\bar{q} = 2$). Profit $\pi_i = p q_i$.

Interior FOCs.

$$a - 3q_1^2 - 8q_1q_2 - 4q_2^2 = 0, \quad a - q_1^2 - 8q_1q_2 - 12q_2^2 = 0.$$

The two equations are not symmetric — no formula like $q_i^* = a/3$.

Existence. $T(q_1, q_2) = (BR_1(q_2), BR_2(q_1))$ is continuous on the compact convex set $K = [0, \bar{q}]^2$.

Brouwer: a Nash equilibrium exists.

3.2 Separating Hyperplane Theorem

Setup. Nonempty **convex** sets $A, B \subset \mathbb{R}^n$ with $A \cap B = \emptyset$.

Theorem. If at least one of A, B is **compact**, then there exist $\mathbf{p} \neq \mathbf{0}$ and $\beta \in \mathbb{R}$ such that

$$\mathbf{p} \cdot \mathbf{a} \leq \beta \leq \mathbf{p} \cdot \mathbf{b} \quad \text{for all } \mathbf{a} \in A, \mathbf{b} \in B.$$

Geometrically, the hyperplane $\mathbf{p} \cdot \mathbf{x} = \beta$ separates A from B . (If only B is compact, swap the labels of A and B .)

Proof and strict separation: [appendix \(for interest only\)](#).


3.2 Separating Hyperplane Theorem — two goods

With $x = (x_1, x_2)$, separation is a line $p_1x_1 + p_2x_2 = \beta$: all of A on one side, all of B on the other.

Two useful special cases.

- **Single point.** $B = \{b_0\}$ is compact. If $b_0 \notin A$ and A is closed and convex, a separator exists even when A is not compact.
- **Strict separation.** If the sets are a positive distance apart, the same $\mathbf{p}$ can separate **strictly**:
 $\mathbf{p} \cdot a < \beta < \mathbf{p} \cdot b$.

Disjoint convex sets separated by a line



3.2 Example — producer theory (numerical illustration)

Netput plan $y = (y_1, y_2)$: input $y_1 \leq 0$, output $y_2 \geq 0$ (second quadrant).

Technology. With $L = -y_1$ units of labor, max output is $f(L) = 2\sqrt{L}$, so


$$Y = \{(y_1, y_2) : y_1 \leq 0, 0 \leq y_2 \leq 2\sqrt{-y_1}\} \quad (\text{free disposal on output}).$$

Prices. Wage $p_1 = 1$, output price $p_2 = 2$. Profit $\pi = p_1 y_1 + p_2 y_2$ (since $y_1 < 0$, the term $p_1 y_1$ is labor cost).

Efficient plan. Use $L = 4$: $y^* = (-4, 4)$. Then $\pi = 1 \cdot (-4) + 2 \cdot 4 = 4$.

Isoprofit line. $y_1 + 2y_2 = 4$, or $y_2 = 2 - \frac{1}{2}y_1$. At $y_1 = 0$: y_2 -axis intercept $y_2 = \pi/p_2 = 2$. Slope $-p_1/p_2 = -\frac{1}{2}$ matches the frontier at y^* (supporting prices).

3.2 Example — producer theory (graph)



3.2 Example — producer theory

Setup. A firm's **production possibility set** $Y \subset \mathbb{R}^n$ is **convex** (and closed in full treatments).

Statement. If Y is convex, the theorem guarantees that every **technically efficient** production plan $y \in Y$ is also a **profit-maximizing** choice for some nonzero price vector $\mathbf{p}$:

$$\mathbf{p} \cdot y \geq \mathbf{p} \cdot y' \quad \text{for all } y' \in Y.$$

Separation picture. Let B be outputs **strictly better than** y (more of some good, no less of any). If y is efficient, $Y \cap B = \emptyset$. A separating hyperplane gives supporting prices $\mathbf{p}$ at y .

Proof: [appendix](#) (for interest only).

Linear Algebra Basics

3.3 Vectors in $\mathbb{R}^n$

Vector. An ordered list of n real numbers. We write a **column vector** as

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} \in \mathbb{R}^n.$$

Its **transpose** $\mathbf{x}^T = (x_1, \dots, x_n)$ is a **row vector**.


Linear combination. Given vectors $\mathbf{v}_1, \dots, \mathbf{v}_k \in \mathbb{R}^n$ and scalars $c_1, \dots, c_k$,

$$c_1 \mathbf{v}_1 + c_2 \mathbf{v}_2 + \dots + c_k \mathbf{v}_k$$

is a **linear combination** of the $\mathbf{v}_j$.

Span. All linear combinations of $\mathbf{v}_1, \dots, \mathbf{v}_k$ form the **span** of those vectors.

3.3 Span in the plane (two vectors)



3.3 Matrix–vector product — columns of A

Setup. $A \in \mathbb{R}^{m \times n}$ has columns $\mathbf{a}_1, \dots, \mathbf{a}_n \in \mathbb{R}^m$ and $\mathbf{x} \in \mathbb{R}^n$ is a column vector.

Column view. $A\mathbf{x}$ is a linear combination of the **columns** of A :

$$A\mathbf{x} = x_1\mathbf{a}_1 + x_2\mathbf{a}_2 + \cdots + x_n\mathbf{a}_n \in \mathbb{R}^m.$$

The weights are the entries of $\mathbf{x}$.

Example. $A = \begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix}$, $\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$:

$$A\mathbf{x} = x_1 \begin{pmatrix} 1 \\ 3 \\ 5 \end{pmatrix} + x_2 \begin{pmatrix} 2 \\ 4 \\ 6 \end{pmatrix}.$$

3.3 Row vector times matrix — rows of A

Setup. $A \in \mathbb{R}^{m \times n}$, $\mathbf{x} \in \mathbb{R}^n$ (column), so $A\mathbf{x} \in \mathbb{R}^m$.

Row view (transpose both sides). $(A\mathbf{x})^T = \mathbf{x}^T A^T$:

$$\mathbf{x}^T A^T = x_1 \mathbf{a}_1^T + x_2 \mathbf{a}_2^T + \cdots + x_n \mathbf{a}_n^T \in \mathbb{R}^{1 \times m}.$$

Dimensions: $(1 \times n)(n \times m) = (1 \times m)$. **Note:** $\mathbf{x}^T A$ does **not** match when $A \in \mathbb{R}^{m \times n}$ with $m \neq n$.

Example. From the previous slide,

$$\mathbf{x}^T = (x_1, x_2), \quad A^T = \begin{pmatrix} 1 & 3 & 5 \\ 2 & 4 & 6 \end{pmatrix}.$$

Then

$$\mathbf{x}^T A^T = x_1(1, 3, 5) + x_2(2, 4, 6).$$

3.3 Two views of the same product

Same numbers, two interpretations. If $A \in \mathbb{R}^{m \times n}$ and $\mathbf{x} \in \mathbb{R}^n$:

- **Column view:** $A\mathbf{x} = \sum_{j=1}^n x_j \mathbf{a}_j \in \mathbb{R}^m$.
- **Entry view:** $(A\mathbf{x})_i = \mathbf{r}_i^T A\mathbf{x} = \sum_{j=1}^n a_{ij}x_j$ (dot product of row i of A with $\mathbf{x}$), where $\mathbf{r}_i \in \mathbb{R}^m$ has i -th entry 1 and all other entries 0 ($\mathbf{r}_i = \mathbf{e}_i$).

Transpose link. $(A\mathbf{x})^T = \mathbf{x}^T A^T$: a **row vector** on the left combines **rows of A^T** (i.e. transposes of columns of A).

3.3 Matrix multiplication AB — example

Example. $A = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}$, $B = \begin{pmatrix} 5 & 6 \\ 7 & 8 \end{pmatrix}$.

Column view (both columns of AB).

$$(AB)\mathbf{e}_1 = A \begin{pmatrix} 5 \\ 7 \end{pmatrix} = 5 \begin{pmatrix} 1 \\ 3 \end{pmatrix} + 7 \begin{pmatrix} 2 \\ 4 \end{pmatrix} = \begin{pmatrix} 19 \\ 43 \end{pmatrix},$$

$$(AB)\mathbf{e}_2 = A \begin{pmatrix} 6 \\ 8 \end{pmatrix} = 6 \begin{pmatrix} 1 \\ 3 \end{pmatrix} + 8 \begin{pmatrix} 2 \\ 4 \end{pmatrix} = \begin{pmatrix} 22 \\ 50 \end{pmatrix}.$$

So $AB = \begin{pmatrix} 19 & 22 \\ 43 & 50 \end{pmatrix}$.

Row view (practice). Find rows 1 and 2 of AB using $(1,2)B$ and $(3,4)B$.

Note (skip in class). $(1,2)B = (19,22)$ and $(3,4)B = (43,50)$ — same matrix as above.

3.3 Matrix multiplication AB

Setup. $A \in \mathbb{R}^{m \times n}$, $B \in \mathbb{R}^{n \times p}$, so $AB \in \mathbb{R}^{m \times p}$.

Column view. Column j of AB is A times column j of B :

$$(AB)\mathbf{e}_j = A(B\mathbf{e}_j) = A\mathbf{b}_j = b_{1j}\mathbf{a}_1 + b_{2j}\mathbf{a}_2 + \cdots + b_{nj}\mathbf{a}_n.$$

Each column of AB is a linear combination of **columns of A** (weights from column j of B).

Row view. Row i of AB is row i of A times B :

$$\mathbf{r}_i^T(AB) = (\mathbf{r}_i^T A)B = a_{i1}\mathbf{s}_1^T + \cdots + a_{in}\mathbf{s}_n^T,$$

where $\mathbf{r}_i \in \mathbb{R}^m$ has i -th entry 1 and all other entries 0 ($\mathbf{r}_i = \mathbf{e}_i$), so $\mathbf{r}_i^T$ selects row i ; and $\mathbf{s}_k^T$ is row k of B : a linear combination of **rows of B** .

3.3 Basis

Idea. A **basis** of $\mathbb{R}^n$ is a small set of vectors from which every vector in $\mathbb{R}^n$ can be built as a **linear combination** in exactly one way.

Standard basis of $\mathbb{R}^n$. For $j = 1, \dots, n$, define $\mathbf{e}_j$ to have a 1 in entry j and 0 elsewhere:

$$\mathbf{e}_1 = \begin{pmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix}, \quad \mathbf{e}_2 = \begin{pmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{pmatrix}, \quad \dots, \quad \mathbf{e}_n = \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{pmatrix}.$$

Coordinates. Every $\mathbf{x} \in \mathbb{R}^n$ can be written uniquely as

$$\mathbf{x} = x_1\mathbf{e}_1 + x_2\mathbf{e}_2 + \dots + x_n\mathbf{e}_n,$$

where $x_1, \dots, x_n$ are the **coordinates** (entries) of $\mathbf{x}$.

3.3 Basis — example ($\mathbb{R}^2$)

Key point. A basis is **not unique**, but for a **fixed** basis every vector has a **unique** representation as a linear combination.

Two bases of $\mathbb{R}^2$. Standard: $\mathbf{e}_1 = (1, 0)$, $\mathbf{e}_2 = (0, 1)$. Alternative: $\mathbf{v}_1 = (1, 0)$, $\mathbf{v}_2 = (1, 1)$.

Same vector, two representations. For $\mathbf{b} = (2, 3)$:

$$\mathbf{b} = 2\mathbf{e}_1 + 3\mathbf{e}_2 \quad \text{and} \quad \mathbf{b} = (-1)\mathbf{v}_1 + 3\mathbf{v}_2.$$

Different bases give different weights $x_1, \dots, x_n$, but within one basis the weights are **unique**.

Requirement for a basis $\{\mathbf{a}_1, \dots, \mathbf{a}_n\}$. Every $\mathbf{b} \in \mathbb{R}^n$ can be written as $\mathbf{b} = x_1\mathbf{a}_1 + \dots + x_n\mathbf{a}_n$ in exactly one way.

3.3 Basis — existence and uniqueness

Requirement for a basis $\{\mathbf{a}_1, \dots, \mathbf{a}_n\}$. Every $\mathbf{b} \in \mathbb{R}^n$ can be written as $\mathbf{b} = x_1\mathbf{a}_1 + \dots + x_n\mathbf{a}_n$ in exactly one way.

Two parts.

- **Existence:** every $\mathbf{b} \in \mathbb{R}^n$ can be written as some linear combination of $\mathbf{a}_1, \dots, \mathbf{a}_n$.
- **Uniqueness:** there is **only one** choice of weights $x_1, \dots, x_n$.

For $A\mathbf{x} = \mathbf{b}$ **with** $A \in \mathbb{R}^{m \times n}$. Then $\mathbf{x} \in \mathbb{R}^n$, $\mathbf{b} \in \mathbb{R}^m$, and $A = [\mathbf{a}_1 \ \dots \ \mathbf{a}_n]$ with $\mathbf{a}_j \in \mathbb{R}^m$ gives

$$A\mathbf{x} = \mathbf{b} \iff \mathbf{b} = x_1\mathbf{a}_1 + \dots + x_n\mathbf{a}_n.$$

3.3 Existence — column space

Setup. $A \in \mathbb{R}^{m \times n}$, $\mathbf{x} \in \mathbb{R}^n$, $\mathbf{b} \in \mathbb{R}^m$, $A = [\mathbf{a}_1 \cdots \mathbf{a}_n]$ with $\mathbf{a}_j \in \mathbb{R}^m$.

Column space.

$$\text{Col}(A) = \text{span}\{\mathbf{a}_1, \dots, \mathbf{a}_n\} \subseteq \mathbb{R}^m.$$

Existence of a solution. $A\mathbf{x} = \mathbf{b}$ has **at least one** solution $\iff \mathbf{b} \in \text{Col}(A)$.

Meaning. $\mathbf{b} \in \text{Col}(A)$ means $\mathbf{b} = x_1\mathbf{a}_1 + \cdots + x_n\mathbf{a}_n$ for some $\mathbf{x}$ — the **existence** part of the basis requirement, but only for this one $\mathbf{b} \in \mathbb{R}^m$ (and only when $\text{Col}(A) = \mathbb{R}^m$ does it hold for **all** $\mathbf{b}$).

3.3 Uniqueness — linearly independent columns

Setup. $A \in \mathbb{R}^{m \times n}$, $\mathbf{x} \in \mathbb{R}^n$, $\mathbf{b} \in \mathbb{R}^m$.

Linearly independent. Vectors $\mathbf{v}_1, \dots, \mathbf{v}_k$ are **linearly independent** if

$$c_1 \mathbf{v}_1 + c_2 \mathbf{v}_2 + \dots + c_k \mathbf{v}_k = \mathbf{0} \implies c_1 = c_2 = \dots = c_k = 0.$$

Linearly dependent: not independent — some nontrivial combination equals $\mathbf{0}$.

Uniqueness of a solution. If $A\mathbf{x} = \mathbf{b}$ has a solution, it is **unique** iff the columns $\mathbf{a}_1, \dots, \mathbf{a}_n$ are **linearly independent**.

3.3 Linear independence $\Leftrightarrow$ unique solution

Claim. Fix $\mathbf{b} \in \mathbb{R}^m$. If $A\mathbf{x} = \mathbf{b}$ has **at least one** solution, then it has a **unique** solution iff $\mathbf{a}_1, \dots, \mathbf{a}_n$ are linearly independent.

$(\Rightarrow)$ **Independent $\Rightarrow$ unique.** If $A\mathbf{x} = \mathbf{b}$ and $A\mathbf{y} = \mathbf{b}$, then $A(\mathbf{x} - \mathbf{y}) = \mathbf{0}$, i.e.

$$(x_1 - y_1)\mathbf{a}_1 + \dots + (x_n - y_n)\mathbf{a}_n = \mathbf{0}.$$

Independence gives $x_j - y_j = 0$ for all j , so $\mathbf{x} = \mathbf{y}$.

$(\Leftarrow)$ **Dependent $\Rightarrow$ not unique.** If the columns are dependent, some $\mathbf{d} \neq \mathbf{0}$ satisfies $A\mathbf{d} = \mathbf{0}$. If $\mathbf{x}$ solves $A\mathbf{x} = \mathbf{b}$, then so does $\mathbf{x} + \mathbf{d}$:

$$A(\mathbf{x} + \mathbf{d}) = A\mathbf{x} + A\mathbf{d} = \mathbf{b} + \mathbf{0} = \mathbf{b},$$

and $\mathbf{x} + \mathbf{d} \neq \mathbf{x}$. So the solution is **not** unique.

3.3 Linear independence — examples ($\mathbb{R}^2$)

Independent (bases).

- Standard: $\mathbf{e}_1 = (1, 0)$, $\mathbf{e}_2 = (0, 1)$.
- Alternative: $\mathbf{v}_1 = (1, 0)$, $\mathbf{v}_2 = (1, 1)$.

For each pair, $c_1\mathbf{w}_1 + c_2\mathbf{w}_2 = \mathbf{0}$ implies $c_1 = c_2 = 0$.

Dependent (too many vectors). $\mathbf{e}_1 = (1, 0)$, $\mathbf{e}_2 = (0, 1)$, and $\mathbf{u} = (1, 1)$:

$$\mathbf{u} = \mathbf{e}_1 + \mathbf{e}_2 \quad \Longleftrightarrow \quad \mathbf{e}_1 + \mathbf{e}_2 - \mathbf{u} = \mathbf{0}.$$

Coefficients $(1, 1, -1)$ are not all zero, so these three vectors are **linearly dependent**. A basis of $\mathbb{R}^2$ has exactly 2 vectors, not 3.

3.3 Linear dependence — columns example

Example (columns of A). $A = \begin{pmatrix} 1 & 2 \\ 2 & 4 \end{pmatrix}$, columns $\mathbf{a}_1 = (1, 2)^T$, $\mathbf{a}_2 = (2, 4)^T$.

Same direction. $\mathbf{a}_2 = 2\mathbf{a}_1$, so the two columns lie on one line through the origin.

Nontrivial combination to 0.

$$2\mathbf{a}_1 - \mathbf{a}_2 = \mathbf{0} \iff 2 \begin{pmatrix} 1 \\ 2 \end{pmatrix} - \begin{pmatrix} 2 \\ 4 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}.$$

Coefficients $(2, -1)$ are not both zero, so $\mathbf{a}_1, \mathbf{a}_2$ are **linearly dependent** (only one independent direction in $\text{Col}(A)$).

3.3 Solving $A\mathbf{x} = \mathbf{b}$ — example

Example. $A = \begin{pmatrix} 1 & 2 \\ 2 & 4 \end{pmatrix}$, so $\text{Col}(A) = \text{span}\{(1, 2)\}$ (a line in $\mathbb{R}^2$).

Solvable. $\mathbf{b} = \begin{pmatrix} 3 \\ 6 \end{pmatrix} = 3(1, 2)$, so $A\mathbf{x} = \mathbf{b}$ has solutions (e.g. $\mathbf{x} = (1, 1)^T$; not unique).

Not solvable (existence fails). $\mathbf{b} = \begin{pmatrix} 1 \\ 0 \end{pmatrix} \notin \text{Col}(A)$, so $A\mathbf{x} = \mathbf{b}$ has **no** solution.

Dependent columns (uniqueness fails). $\mathbf{a}_2 = 2\mathbf{a}_1$, so the columns are **linearly dependent**. When a solution exists, it is not unique (e.g. both $\mathbf{x} = (3, 0)^T$ and $\mathbf{x} = (1, 1)^T$ give $\mathbf{b} = (3, 6)^T$).

3.3 Full basis requirement — only when $m = n$

Goal. When does **every** $\mathbf{b} \in \mathbb{R}^m$ satisfy $\mathbf{b} = x_1\mathbf{a}_1 + \cdots + x_n\mathbf{a}_n$ in **exactly one way**?

Claim. For $A \in \mathbb{R}^{m \times n}$, this happens **only if** $m = n$.

Proof sketch.

- **Uniqueness** $\Rightarrow$ columns independent. In $\mathbb{R}^m$ at most m vectors can be independent, so $n \leq m$.
- **Existence for all $\mathbf{b}$** $\Rightarrow \text{Col}(A) = \mathbb{R}^m$. With only n columns, $\dim(\text{Col}(A)) \leq n$, so $n \geq m$.
- Therefore $m = n$.

If $n < m$, columns cannot span $\mathbb{R}^m$; if $n > m$, n columns in $\mathbb{R}^m$ cannot be independent.

3.3 Square case — basis = existence + uniqueness

Now $A \in \mathbb{R}^{n \times n}$, $\mathbf{b} \in \mathbb{R}^n$. Requirement for a basis $\{\mathbf{a}_1, \dots, \mathbf{a}_n\}$: every $\mathbf{b} \in \mathbb{R}^n$ is $\mathbf{b} = x_1\mathbf{a}_1 + \dots + x_n\mathbf{a}_n$ in exactly one way.

Equivalent to $A\mathbf{x} = \mathbf{b}$.

$$A\mathbf{x} = \mathbf{b} \iff \mathbf{b} = x_1\mathbf{a}_1 + \dots + x_n\mathbf{a}_n.$$

Columns form a basis of $\mathbb{R}^n$ $\iff$ existence and uniqueness for **every $\mathbf{b}$** $\iff A\mathbf{x} = \mathbf{b}$ has a **unique** solution for every $\mathbf{b} \in \mathbb{R}^n$.

3.4 Two views of a matrix A

View 1 — $\mathbf{x}$ as weights. Write $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n]$. Then

$$A\mathbf{x} = x_1\mathbf{a}_1 + x_2\mathbf{a}_2 + \cdots + x_n\mathbf{a}_n.$$

The entries of $\mathbf{x}$ are the **weights** on the columns of A .

View 2 — A as a linear transformation. Define $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$ by $T(\mathbf{x}) = A\mathbf{x}$. The matrix A sends each input vector $\mathbf{x}$ to an output vector in $\mathbb{R}^m$; for square A , $|\det(A)|$ is the area-scaling factor (Section 3.5).

3.4 Linear transformation — definition

Linear transformation. A map $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is **linear** if for all $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ and all scalars $c \in \mathbb{R}$:

$$T(\mathbf{x} + \mathbf{y}) = T(\mathbf{x}) + T(\mathbf{y}), \quad T(c\mathbf{x}) = cT(\mathbf{x}).$$

Equivalently, for all $\mathbf{x}, \mathbf{y}$ and c :

$$T(c\mathbf{x} + \mathbf{y}) = cT(\mathbf{x}) + T(\mathbf{y}).$$

Examples. Rotation, reflection, projection, and shear in the plane are linear. A map with a constant term, e.g. $T(\mathbf{x}) = A\mathbf{x} + \mathbf{b}$ with $\mathbf{b} \neq \mathbf{0}$, is **not** linear.

Matrix form. Every linear T is $T(\mathbf{x}) = A\mathbf{x}$ for some matrix A (general recipe on the next slides). For $A \in \mathbb{R}^{2 \times 2}$, $|\det(A)|$ is the area-scaling factor (Section 3.5).

3.4 Linear maps — decompose, then combine

Benefit. A linear map $T : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ may be hard to visualize on an **arbitrary** input $\mathbf{v}$.

But you often know T on a few vectors. Suppose you already know $T(\mathbf{x})$ and $T(\mathbf{y})$, and $\mathbf{v}$ can be decomposed as

$$\mathbf{v} = a\mathbf{x} + b\mathbf{y}$$

for some scalars $a, b \in \mathbb{R}$. Then by linearity,

$$T(\mathbf{v}) = T(a\mathbf{x} + b\mathbf{y}) = aT(\mathbf{x}) + bT(\mathbf{y}).$$

Decompose $\mathbf{v}$ into known pieces, transform each piece, then combine.

Standard basis recipe ($\mathbb{R}^2$). Take $\mathbf{x} = \mathbf{e}_1$, $\mathbf{y} = \mathbf{e}_2$, so $\mathbf{v} = v_1\mathbf{e}_1 + v_2\mathbf{e}_2$. If you know $T(\mathbf{e}_1)$ and $T(\mathbf{e}_2)$, then $T(\mathbf{v}) = v_1 T(\mathbf{e}_1) + v_2 T(\mathbf{e}_2)$. Stack $T(\mathbf{e}_1)$ and $T(\mathbf{e}_2)$ as columns to get the matrix of T (next slide).

3.4 Decompose, transform, combine — notation

Recipe. For $\mathbf{v} = v_1\mathbf{e}_1 + v_2\mathbf{e}_2$ and linear T :

1. **Decompose** the input:

$$\mathbf{v} = v_1\mathbf{e}_1 + v_2\mathbf{e}_2.$$

2. **Transform** each piece (use known values of T on $\mathbf{e}_1, \mathbf{e}_2$):

$$v_1 T(\mathbf{e}_1), \quad v_2 T(\mathbf{e}_2).$$

3. **Combine** the outputs:

$$T(\mathbf{v}) = v_1 T(\mathbf{e}_1) + v_2 T(\mathbf{e}_2) = A\mathbf{v}, \quad A = \begin{bmatrix} T(\mathbf{e}_1) & T(\mathbf{e}_2) \end{bmatrix}.$$

3.4 Example — shear

Map. Horizontal shear: tilt the unit square into a parallelogram, $(x_1, x_2) \mapsto (x_1 + x_2, x_2)$.



3.4 Example — shear (matrix)

Find the matrix. Apply T to $\mathbf{e}_1$ and $\mathbf{e}_2$, then stack the outputs as columns:

$$T(\mathbf{e}_1) = \mathbf{e}_1 = \begin{pmatrix} 1 \\ 0 \end{pmatrix}, \quad T(\mathbf{e}_2) = \mathbf{e}_1 + \mathbf{e}_2 = \begin{pmatrix} 1 \\ 1 \end{pmatrix} \quad \Rightarrow \quad A = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}.$$

Check: $A\mathbf{e}_1 = \mathbf{e}_1$ (bottom edge fixed); $A\mathbf{e}_2$ tilts the vertical side.

3.4 Example — shear (decompose, transform, combine)

Take $\mathbf{v} = 2\mathbf{e}_1 + 3\mathbf{e}_2 = \begin{pmatrix} 2 \\ 3 \end{pmatrix}$.

1. **Decompose:** $\mathbf{v} = 2\mathbf{e}_1 + 3\mathbf{e}_2$.


2. **Transform:** $2T(\mathbf{e}_1) = 2 \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \begin{pmatrix} 2 \\ 0 \end{pmatrix}, \quad 3T(\mathbf{e}_2) = 3 \begin{pmatrix} 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 3 \\ 3 \end{pmatrix}.$

3. **Combine:**

$$T(\mathbf{v}) = 2T(\mathbf{e}_1) + 3T(\mathbf{e}_2) = \begin{pmatrix} 2 \\ 0 \end{pmatrix} + \begin{pmatrix} 3 \\ 3 \end{pmatrix} = \begin{pmatrix} 5 \\ 3 \end{pmatrix} = A\mathbf{v}.$$

3.4 Example — 90° rotation (clockwise)

Map. Rotate every vector 90° clockwise: $(x_1, x_2) \mapsto (x_2, -x_1)$.



3.4 Example — rotation (matrix)

Find the matrix. Apply T to $\mathbf{e}_1$ and $\mathbf{e}_2$, then stack the outputs as columns:

$$T(\mathbf{e}_1) = \begin{pmatrix} 0 \\ -1 \end{pmatrix}, \quad T(\mathbf{e}_2) = \begin{pmatrix} 1 \\ 0 \end{pmatrix} \implies A = \begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix}.$$

Check: $A\mathbf{e}_1$ points down; $A\mathbf{e}_2$ points right (see previous graph).

3.4 Example — rotation (decompose, transform, combine)

Same recipe with $\mathbf{v} = 2\mathbf{e}_1 + 3\mathbf{e}_2 = \begin{pmatrix} 2 \\ 3 \end{pmatrix}$.

1. **Decompose:** $\mathbf{v} = 2\mathbf{e}_1 + 3\mathbf{e}_2$.

2. **Transform:** $2T(\mathbf{e}_1) = 2 \begin{pmatrix} 0 \\ -1 \end{pmatrix} = \begin{pmatrix} 0 \\ -2 \end{pmatrix}, \quad 3T(\mathbf{e}_2) = 3 \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \begin{pmatrix} 3 \\ 0 \end{pmatrix}.$

3. **Combine:**

$$T(\mathbf{v}) = 2T(\mathbf{e}_1) + 3T(\mathbf{e}_2) = \begin{pmatrix} 0 \\ -2 \end{pmatrix} + \begin{pmatrix} 3 \\ 0 \end{pmatrix} = \begin{pmatrix} 3 \\ -2 \end{pmatrix} = A\mathbf{v}.$$

3.4 Every linear map is a matrix

Claim. If $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is linear, then there is a unique matrix $A \in \mathbb{R}^{m \times n}$ such that $T(\mathbf{x}) = A\mathbf{x}$ for all $\mathbf{x}$.

Construction. Let $\mathbf{e}_1, \dots, \mathbf{e}_n$ be the standard basis of $\mathbb{R}^n$. Form A by stacking the outputs as columns:

$$A = [T(\mathbf{e}_1) \quad T(\mathbf{e}_2) \quad \cdots \quad T(\mathbf{e}_n)].$$

Why it works. Any $\mathbf{x} = \sum_{j=1}^n x_j \mathbf{e}_j$. By linearity,

$$T(\mathbf{x}) = \sum_{j=1}^n x_j T(\mathbf{e}_j) = \sum_{j=1}^n x_j \mathbf{a}_j = A\mathbf{x}.$$

Column j of A is exactly where $\mathbf{e}_j$ is sent.

3.5 Determinant — area scaling (2×2)

Setup. $A = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$, columns $\mathbf{a}_1 = (a, c)^T$, $\mathbf{a}_2 = (b, d)^T$. View A as the linear map $T(\mathbf{x}) = A\mathbf{x}$ (shear, rotation, etc. on previous slides).

Where the standard basis goes.

$$A \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \begin{pmatrix} a \\ c \end{pmatrix} = \mathbf{a}_1, \quad A \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \begin{pmatrix} b \\ d \end{pmatrix} = \mathbf{a}_2.$$


So $(1, 0)^T$ is sent to $(a, c)^T$, and $(0, 1)^T$ to $(b, d)^T$.

Area conversion factor. The unit square (spanned by $(1, 0)^T$ and $(0, 1)^T$, area = 1) is mapped by A to the parallelogram spanned by $\mathbf{a}_1$ and $\mathbf{a}_2$. The **determinant** is exactly the (signed) area of that parallelogram:

$$\det(A) = ad - bc.$$

So $\det(A)$ is the **area conversion factor**: every region's area is multiplied by $|\det(A)|$ under $\mathbf{x} \mapsto A\mathbf{x}$.

3.5 Determinant — why $ad - bc$ is the area



Box area $(a + b)(c + d)$ minus purple (bc) , red $(ac/2)$, teal $(bd/2)$ pieces leaves yellow parallelogram area

$$ad - bc = \det(A).$$

3.5 Determinant = 0 — example

Example. $A = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}$, so $a = b = c = d = 1$ and

$$\det(A) = ad - bc = 1 \cdot 1 - 1 \cdot 1 = 0.$$

Standard basis collapses to a line.

$$A \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \end{pmatrix}, \quad A \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \end{pmatrix}.$$

Both $(1, 0)^T$ and $(0, 1)^T$ are sent to the **same** point on one line through the origin. The unit square (area 1) is squashed to a **1-dimensional** set — area 0.

So (square case):

$$\det(A) = 0 \iff \text{columns linearly dependent.}$$

Here $\mathbf{a}_2 = \mathbf{a}_1 = (1, 1)^T$, so the columns are dependent. **Next:** when $\det(A) = 0$, some input direction is **squeezed** to **0**.

3.5 Null space — directions squeezed to 0

Setup ($A \in \mathbb{R}^{2 \times 2}$). View A as a linear transformation $T : \mathbb{R}^2 \rightarrow \mathbb{R}^2$, $T(\mathbf{x}) = A\mathbf{x}$.

Null space / kernel (subspace of $\mathbb{R}^2$):

$$\text{Null}(A) = \ker(T) = \{\mathbf{x} \in \mathbb{R}^2 : A\mathbf{x} = \mathbf{0}\}.$$

Vectors in $\text{Null}(A)$ are mapped to $\mathbf{0}$ in the output.

Intuition. If $\mathbf{x} \neq \mathbf{0}$ lies in $\text{Null}(A)$, the whole line $\text{span}\{\mathbf{x}\} \subset \mathbb{R}^2$ is sent to the single point $\mathbf{0} \in \mathbb{R}^2$ — that **input direction is squeezed** (lost in the output).

Example ($\det A = 0$). $A = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}$: $A(1, -1)^T = \mathbf{0}$, so $\text{Null}(A) = \text{span}\{(1, -1)^T\}$ — one direction squeezed.

Next: how many directions squeeze vs. persist?

3.5 Rank–nullity — squeezed vs. persist (2×2)

Two fates ($A \in \mathbb{R}^{2 \times 2}$). Under $T(\mathbf{x}) = A\mathbf{x}$, each input direction either:

1. **Squeezed to 0** $\ker(T) = \text{Null}(A)$ (previous slide); dimension = nullity(A).
2. **Persist in the output** $\text{Im}(T) = \text{Col}(A) = \{A\mathbf{x} : \mathbf{x} \in \mathbb{R}^2\}$; dimension = rank(A) — independent directions that **survive**.

Rank–nullity. $\text{rank}(A) + \text{nullity}(A) = 2$.

Example ($\det A = 0$). Collapse matrix from previous slides: nullity = 1, rank = 1 (outputs fill $\text{span}\{(1, 1)^T\}$).

Example ($\det A \neq 0$). Shear or rotation (Section 3.4): nullity = 0, rank = 2.

3.5 After a squeeze — what survives? (2×2)

Recall. $A = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}$ maps the whole plane onto one output line (through $(1, 1)^T$). Direction $(1, -1)^T$ is squeezed to $\mathbf{0}$.

Question. Given an input $\mathbf{v}$, what is $A\mathbf{v}$?

Key idea (in words). Think of $\mathbf{v}$ as two moves: one along the **squeezed** direction (the map **forgets** it) and one along the surviving direction $(1, 1)^T$ (the map **remembers** only this). So $A\mathbf{v}$ depends on how much of $\mathbf{v}$ points along the surviving direction — not on the part that collapses.

Example. $(3, 1)^T = 2(1, 1)^T + 1(1, -1)^T$: the $(1, -1)^T$ part is lost, the $2(1, 1)^T$ part survives, so $A\mathbf{v} = 2(2, 2)^T = (4, 4)^T$.

Next: the same squeeze/persist picture in general $m \times n$.

3.5 Rank–nullity — general case **for interest only**

Setup. Let $A \in \mathbb{R}^{m \times n}$, $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$, $T(\mathbf{x}) = A\mathbf{x}$.

Kernel and image. $\ker(T) = \text{Null}(A) = \{\mathbf{x} \in \mathbb{R}^n : A\mathbf{x} = \mathbf{0}\} \subseteq \mathbb{R}^n$;
 $\text{Im}(T) = \text{Col}(A) = \{A\mathbf{x} : \mathbf{x} \in \mathbb{R}^n\} \subseteq \mathbb{R}^m$.

Rank and nullity. $\text{rank}(A) = \dim(\text{Im } T)$, $\text{nullity}(A) = \dim(\ker T)$.

Rank–nullity theorem.

$$\text{rank}(A) + \text{nullity}(A) = n \quad (n \text{ columns} = \dim(\mathbb{R}^n)).$$

Same intuition. Of the n domain dimensions, $\text{nullity}(A)$ are **squeezed** to $\mathbf{0} \in \mathbb{R}^m$; $\text{rank}(A)$ independent directions **persist** in $\text{Im}(T) \subseteq \mathbb{R}^m$. When $n = m$, $\det(A) \neq 0 \Leftrightarrow \text{nullity} = 0 \Leftrightarrow \text{rank} = n$ (same as invertible / basis on the next slide).

3.5 Invertible matrices — column picture

Recall ($A \in \mathbb{R}^{n \times n}$). $Ax = b \iff b = x_1 a_1 + \dots + x_n a_n$ (§3.3).

Basis = two requirements (§3.3). Columns $a_1, \dots, a_n$ form a **basis** of $\mathbb{R}^n$ iff:

- (1) **Uniqueness:** at most one choice of weights $x_1, \dots, x_n$ (columns **linearly independent**).
- (2) **Solvability:** every $b \in \mathbb{R}^n$ is such a combination ($\text{Col}(A) = \mathbb{R}^n$).

For square A , (1) $\Leftrightarrow$ (2); together they give a **unique** solution for **every** b , and A^{-1} exists.

If (1) or (2) fails (not a basis):

- (2) fails: $\text{Col}(A) \subsetneq \mathbb{R}^n$ — some b have **no** solution.
- (1) fails (columns **linearly dependent**): $\exists d \neq 0$ with $Ad = 0$ — if a solution exists, $x + d$ is another (not unique).

Link to geometry (this section). $\det(A) = 0 \Leftrightarrow$ a squeeze $\Leftrightarrow$ not invertible; $\det(A) \neq 0 \Leftrightarrow$ nullity = 0 $\Leftrightarrow$ basis $\Leftrightarrow$ invertible.

3.5 Invertible matrices — equivalent conditions

Let $A \in \mathbb{R}^{n \times n}$. All eight statements are **equivalent**.

I. Column picture (§3.3)

1. Columns **linearly independent** — **uniqueness**.
2. $\text{Col}(A) = \mathbb{R}^n$ — **solvability for all $\mathbf{b}$** .
3. Columns form a **basis** — (1) and (2) [square: $(1) \Leftrightarrow (2)$].

II. Solving $A\mathbf{x} = \mathbf{b}$

4. $A\mathbf{x} = \mathbf{0}$ only $\mathbf{x} = \mathbf{0}$.
5. For every $\mathbf{b}$, **unique** solution $\mathbf{x} = A^{-1}\mathbf{b}$.

III. Matrix / geometry

6. A **invertible** ($AA^{-1} = I_n$).
7. $\det(A) \neq 0$ (no squeeze).
8. Rows form a **basis** of $\mathbb{R}^n$.

3.5 Transpose and inverse — rules we use

Transpose (always defined).

$$(A^T)^T = A, \quad (AB)^T = B^T A^T, \quad (Ax)^T = x^T A^T.$$

Why (8)? (2) $\Rightarrow Ax = \mathbf{b}$ for all $\mathbf{b}$. Transpose: $x^T A^T = \mathbf{b}^T$ for all $\mathbf{b}^T$ (row view, §3.3). Rows of $A =$ columns of $A^T \Rightarrow$ rows span $\mathbb{R}^n$; square $\Rightarrow$ basis [(8)].

Inverse (definition). For square $A \in \mathbb{R}^{n \times n}$, A^{-1} satisfies $AA^{-1} = A^{-1}A = I_n$. It exists iff any condition on the previous slide holds (e.g. $\det(A) \neq 0$). Then $Ax = \mathbf{b}$ has unique solution $\mathbf{x} = A^{-1}\mathbf{b}$.

Inverse rules (both factors invertible).

$$(AB)^{-1} = B^{-1}A^{-1}, \quad (A^{-1})^{-1} = A, \quad (A^{-1})^T = (A^T)^{-1}.$$

Proofs: [appendix](#) (for interest only).

3.6 Dot product — algebraic and geometric

For $\mathbf{u}, \mathbf{v} \in \mathbb{R}^n$:

Algebraic definition.

$$\mathbf{u} \cdot \mathbf{v} = \sum_{i=1}^n u_i v_i.$$

Geometric definition. Let θ be the angle between $\mathbf{u}$ and $\mathbf{v}$. Then

$$\mathbf{u} \cdot \mathbf{v} = \|\mathbf{u}\| \|\mathbf{v}\| \cos \theta.$$

Equivalent. These two definitions always give the **same number**.

Video (intuition and proof). <https://www.youtube.com/watch?v=LyGKycYT2v0>

3.6 Linear regression — $\mathbf{y}$ and X

Data (n workers). Worker i has income y_i , age age_i , education educ_i .

Response vector.

$$\mathbf{y} = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix} \in \mathbb{R}^n.$$


Regressor matrix (two columns: age, educ).

$$X = \begin{pmatrix} \text{age}_1 & \text{educ}_1 \\ \vdots & \vdots \\ \text{age}_n & \text{educ}_n \end{pmatrix} \in \mathbb{R}^{n \times 2}, \quad \boldsymbol{\beta} = \begin{pmatrix} \beta_{\text{age}} \\ \beta_{\text{educ}} \end{pmatrix}, \quad \mathbf{y} \approx X\boldsymbol{\beta}.$$

Column j of X is the vector of worker i values for regressor j .

Schematic next ($n = 3$): age and education are the two axes; yellow plane = $\text{Col}(X)$; $\mathbf{e}$ is orthogonal to every column of X .

Yellow plane
=
 $\text{Col}(X)$



3.6 OLS fit — $\mathbf{e}$ orthogonal to every column of X

From the picture ($n = 3$). Yellow plane = $\text{Col}(X)$; fitted $\hat{\mathbf{y}} = X\hat{\boldsymbol{\beta}}$ lies in it. Residual $\mathbf{e} = \mathbf{y} - X\hat{\boldsymbol{\beta}}$ is perpendicular to the plane.

Orthogonal to every column of X . Let $\mathbf{x}_1 = \text{age column}$, $\mathbf{x}_2 = \text{educ column}$. Then

$$\mathbf{x}_1 \cdot \mathbf{e} = 0, \quad \mathbf{x}_2 \cdot \mathbf{e} = 0 \quad \Longleftrightarrow \quad X^T \mathbf{e} = \mathbf{0} \quad \Longleftrightarrow \quad X^T (\mathbf{y} - X\hat{\boldsymbol{\beta}}) = \mathbf{0}.$$

These **normal equations** characterise $\hat{\boldsymbol{\beta}}$.

Solve for $\hat{\boldsymbol{\beta}}$. $X^T X \hat{\boldsymbol{\beta}} = X^T \mathbf{y}$. If the columns of X are independent (full column rank), then

$$\hat{\boldsymbol{\beta}} = (X^T X)^{-1} X^T \mathbf{y}.$$

(Since $\hat{\mathbf{y}} \in \text{Col}(X)$, also $\hat{\mathbf{y}} \cdot \mathbf{e} = 0$.)

3.6 OLS fit — same as least squares

Pythagoras in $\mathbb{R}^n$. Fix $\hat{\beta}$ with $\mathbf{e} = \mathbf{y} - X\hat{\beta}$. For any β ,

$$\mathbf{y} - X\beta = \underbrace{\mathbf{y} - X\hat{\beta}}_{\mathbf{e}} + \underbrace{X(\hat{\beta} - \beta)}_{\in \text{Col}(X)}.$$

The two pieces are orthogonal, so

$$\|\mathbf{y} - X\beta\|^2 = \|\mathbf{y} - X\hat{\beta}\|^2 + \|X(\hat{\beta} - \beta)\|^2 \geq \|\mathbf{y} - X\hat{\beta}\|^2.$$

Conclusion. The smallest squared error occurs when $X(\hat{\beta} - \beta) = \mathbf{0}$, i.e.
 $\hat{\beta} = \arg \min_{\beta} \|\mathbf{y} - X\beta\|^2 = \arg \min_{\beta} \sum_{i=1}^n (y_i - \beta_{\text{age}} \text{age}_i - \beta_{\text{educ}} \text{educ}_i)^2.$

3.6 The projection matrix

Projection matrix P_X (hat matrix). If $X \in \mathbb{R}^{n \times k}$ has full column rank, define

$$P_X = X(X^T X)^{-1} X^T.$$

Then $\hat{\mathbf{y}} = P_X \mathbf{y}$.

Annihilator matrix M_X (residual maker matrix). Define $M_X = I - P_X$. Then $\mathbf{e} = M_X \mathbf{y}$; and $M_X X = \mathbf{0}$ (M_X annihilates $\text{Col}(X)$).

Useful facts. $P_X^T = P_X$, $P_X^2 = P_X$, $M_X X = \mathbf{0}$, and $M_X^T = M_X$, $M_X^2 = M_X$.

Why $P_X^T = P_X$ and $P_X^2 = P_X$: [appendix \(for interest only\)](#).

3.6 Frisch–Waugh–Lovell (FWL) theorem for interest only

Partition regressors. Write $X = [X_1 \mid X_2]$ (e.g. age = X_1 , educ = X_2). Let P_{X_2} be the **projection matrix** onto $\text{Col}(X_2)$ and $M_{X_2} = I - P_{X_2}$ the **annihilator matrix** (residual maker matrix; write P_2 , M_2 for short).

Claim. The OLS coefficient on X_1 in $y \sim X_1 + X_2$ equals the OLS coefficient in

$$M_2 y \sim M_2 X_1$$

(M_2 removes the part of each vector explained by educ).

Reading. “Controlling for educ” = apply the annihilator M_2 to both income and age, then regress.

Proof idea. Write $y = X_1 \hat{\beta}_1 + X_2 \hat{\beta}_2 + e$ with $X^T e = 0$. Since $M_2 X_2 = 0$ and $M_2 e = e$, get $M_2 y = M_2 X_1 \hat{\beta}_1 + e$; also $(M_2 X_1)^T e = X_1^T M_2 e = X_1^T e = 0$, so $\hat{\beta}_1$ solves the partialled normal equations.

3.7 Which basis makes decompose–transform–combine easy?

Recall the recipe. If $\mathbf{v} = a\mathbf{x} + b\mathbf{y}$, then $T(\mathbf{v}) = aT(\mathbf{x}) + bT(\mathbf{y})$ — always true by linearity.

But step 2 can be hard. You must know $T(\mathbf{x})$ and $T(\mathbf{y})$. With the standard basis $(\mathbf{e}_1, \mathbf{e}_2)$, rotation sent $\mathbf{e}_1$ down and $\mathbf{e}_2$ right — both outputs are **new directions** to track.

A proper basis (for this map). Pick directions where T **only scales**, not rotates or shears:

$$T(\mathbf{v}) = \lambda\mathbf{v} \quad \text{for some scalar } \lambda.$$

Then $T(a\mathbf{v}) = a\lambda\mathbf{v}$ — one number λ , same direction. These special directions are **eigenvectors**; λ is the **eigenvalue**.

3.7 Eigenvectors and eigenvalues — definition

Setup. Square matrix $A \in \mathbb{R}^{n \times n}$ (linear map $T(\mathbf{x}) = A\mathbf{x}$).

Eigenvector / eigenvalue. A nonzero vector $\mathbf{v}$ is an **eigenvector** of A if

$$A\mathbf{v} = \lambda\mathbf{v}$$

for some scalar λ (the **eigenvalue**). Geometrically: A stretches or flips $\mathbf{v}$, but keeps its line.

Proper basis. When A has enough independent eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_n$, decompose in **that** basis instead of the standard one.

If $A = A^T$, eigenvalues are real (**appendix**) (**for interest only**).

3.7 Finding eigenvalues and eigenvectors

Traditional recipe (works in any dimension):

1. **Eigenvalues.** Solve the **characteristic equation**

$$\det(A - \lambda I) = 0.$$

2. **Eigenvectors.** For each eigenvalue λ , solve

$$(A - \lambda I)\mathbf{v} = \mathbf{0}.$$

Nontrivial solutions are eigenvectors; the solution set is the **eigenspace** E_λ .

Small n (e.g. 2×2): solve the equations by hand (Together exercise below).

General / high n : row-reduce $(A - \lambda I)$ to **RREF** and read off pivot vs. free columns — **appendix (for interest only)**.

3. **Repeated** λ . If λ has algebraic multiplicity > 1 , $\dim E_\lambda$ equals the number of free variables in step 2.

Together: 2×2 by hand. Take home: $\mathbb{R}^4$ exercise below.

3.7 Exercise: Eigenvectors — 2×2 example Together

Exercise. For $A = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$, find all eigenpairs and compute $A\mathbf{v}$ for $\mathbf{v} = \begin{pmatrix} 3 \\ 1 \end{pmatrix}$.

Solution (eigenvalues). Characteristic equation $\det(A - \lambda I) = 0$:

$$(2 - \lambda)^2 - 1 = \lambda^2 - 4\lambda + 3 = (\lambda - 1)(\lambda - 3) = 0 \quad \Rightarrow \quad \lambda_1 = 3, \lambda_2 = 1.$$

Solution ($\lambda_1 = 3$). Solve $(A - 3I)\mathbf{v} = \mathbf{0}$:

$$\begin{pmatrix} -1 & 1 \\ 1 & -1 \end{pmatrix} \begin{pmatrix} v_1 \\ v_2 \end{pmatrix} = \mathbf{0} \quad \Rightarrow \quad -v_1 + v_2 = 0 \quad \Rightarrow \quad \mathbf{v}_1 = \begin{pmatrix} 1 \\ 1 \end{pmatrix}.$$

Solution ($\lambda_2 = 1$). Solve $(A - I)\mathbf{v} = \mathbf{0}$:

$$\begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} v_1 \\ v_2 \end{pmatrix} = \mathbf{0} \quad \Rightarrow \quad v_1 + v_2 = 0 \quad \Rightarrow \quad \mathbf{v}_2 = \begin{pmatrix} 1 \\ -1 \end{pmatrix}.$$

Check: $A\mathbf{v}_1 = 3\mathbf{v}_1$, $A\mathbf{v}_2 = 1\mathbf{v}_2$.

3.7 Exercise: Eigenvectors — 2×2 example (cont.)

Together

Compute $A\mathbf{v}$ for $\mathbf{v} = \begin{pmatrix} 3 \\ 1 \end{pmatrix}$ (decompose–transform–combine).

Step 1 (decompose). Write $\mathbf{v} = c_1\mathbf{v}_1 + c_2\mathbf{v}_2$:

$$c_1 \begin{pmatrix} 1 \\ 1 \end{pmatrix} + c_2 \begin{pmatrix} 1 \\ -1 \end{pmatrix} = \begin{pmatrix} 3 \\ 1 \end{pmatrix} \Rightarrow c_1 = 2, c_2 = 1.$$

Step 2 (transform). A scales each eigen-direction by its eigenvalue:

$$A\mathbf{v} = 2 \cdot 3\mathbf{v}_1 + 1 \cdot 1\mathbf{v}_2 = \begin{pmatrix} 7 \\ 5 \end{pmatrix}.$$

3.7 Exercise: Eigenvectors — 2×2 example (cont.)

Together

Compute $A^{10}\mathbf{v}$ for $\mathbf{v} = \begin{pmatrix} 3 \\ 1 \end{pmatrix} = 2\mathbf{v}_1 + \mathbf{v}_2$.

On an eigenvector: $A^{10}\mathbf{v}_1 = A^9(A\mathbf{v}_1) = A^9(\lambda_1\mathbf{v}_1) = \lambda_1 A^9\mathbf{v}_1 = \dots = \lambda_1^{10}\mathbf{v}_1$ (and similarly $A^{10}\mathbf{v}_2 = \lambda_2^{10}\mathbf{v}_2$).

Same recipe: multiply each coefficient by λ_i^{10} :

$$A^{10}\mathbf{v} = 2 \cdot 3^{10}\mathbf{v}_1 + 1 \cdot 1^{10}\mathbf{v}_2.$$

Since $3^{10} = 59049$,

$$A^{10}\mathbf{v} = 118098 \begin{pmatrix} 1 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ -1 \end{pmatrix} = \begin{pmatrix} 118099 \\ 118097 \end{pmatrix}.$$

The $\lambda_1 = 3$ term dominates; repeated multiplication amplifies the $(1, 1)$ direction.

3.7 Exercise: Eigenvectors in $\mathbb{R}^4$ (RREF) Take home

Exercise. For

$$A = \begin{pmatrix} 3 & 1 & 0 & 0 \\ 1 & 3 & 0 & 0 \\ 0 & 0 & 3 & 1 \\ 0 & 0 & 1 & 3 \end{pmatrix},$$

find all eigenvalues. For $\lambda = 2$, row-reduce $(A - 2I)$ to RREF and find a basis of E_2 (RREF: [appendix](#)).

Solution (eigenvalues). Two decoupled 2×2 blocks; each gives $(3 - \lambda)^2 - 1 = 0$.

$$\det(A - \lambda I) = ((3 - \lambda)^2 - 1)^2 = (\lambda - 2)^2(\lambda - 4)^2.$$

So $\lambda = 2$ and $\lambda = 4$, each with algebraic multiplicity 2.

3.7 Exercise: Eigenvectors in $\mathbb{R}^4$ (RREF, cont.)

Take home

Solution ($\lambda = 2$). Row-reduce

$$A - 2I = \begin{pmatrix} 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 \end{pmatrix} \longrightarrow \begin{pmatrix} 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix}.$$

Pivot columns 1, 3 $\Rightarrow v_1, v_3$ pivot; **free columns** 2, 4 $\Rightarrow v_2, v_4$ free. Equations: $v_1 + v_2 = 0$, $v_3 + v_4 = 0$.

General solution: $(v_1, v_2, v_3, v_4) = v_2(-1, 1, 0, 0) + v_4(0, 0, -1, 1)$.

Basis of E_2 (Appendix A.6 recipe): $(-1, 1, 0, 0)^\top, (0, 0, -1, 1)^\top$.

3.7 Eigenbasis — especially useful in high dimension

If $\mathbf{v}_1, \dots, \mathbf{v}_n$ is an eigenbasis and $\mathbf{v} = \sum_i c_i \mathbf{v}_i$, then

$$A\mathbf{v} = \sum_i c_i \lambda_i \mathbf{v}_i, \quad A^k \mathbf{v} = \sum_i c_i \lambda_i^k \mathbf{v}_i.$$

In an eigenbasis, A is just **separate scalings** — easy to describe and compute when n is large.

Not every matrix has a full real eigenbasis (e.g. plane rotation); **appendix (for interest only)**. SVD/QR play a similar role when it does not.

3.7 Applications — PCA and Markov chains **for interest only**

for interest only You may need to understand some of these when you pursue macro, empirical economics, or IO.

Extremely useful beyond the 2×2 example. Many applied problems reduce to: find directions where a matrix acts simply.

- **PCA (principal components).** Eigenvectors of the **covariance matrix** are the main directions of variation in data; largest eigenvalues pick the most important components for dimension reduction (Section 3.8).
- **Markov transition matrices.** For a stochastic matrix P (rows sum to 1), the **long-run distribution** is an eigenvector with $\lambda = 1$; other eigenvalues control how fast $P^k \mathbf{v}$ converges to that steady state.

Pattern. Decompose a state or data vector in special directions; each direction evolves by a single number (λ_i).

3.7 Applications — ODEs for interest only

Ordinary differential equations (ODEs) — eigenvalues / eigenvectors show up in:

- Linear systems $\dot{x} = Ax$
- Stability of equilibria and phase-plane dynamics
- Continuous-time macro and growth models
- Linearized dynamics around steady states
- Compartment models (e.g. epidemiology, population dynamics)

3.7 Decompositions — QR, eigen, SVD for interest only

Different decompositions, same spirit: rewrite a matrix so that the hard part is **diagonal**, **triangular**, or **orthogonal**.

- **QR decomposition** ($A = QR$). Q has orthonormal columns spanning $\text{Col}(A)$; R is upper triangular. Used to pick **linearly independent columns**, solve least squares stably, and build regression numerically.
- **Spectral / eigen decomposition** ($A = V\Lambda V^{-1}$). When A has a full eigenbasis, V lists eigenvectors and Λ lists eigenvalues — matrix powers and dynamics become coefficient-wise.
- **Singular Value Decomposition (SVD)** ($A = U\Sigma V^T$). Always exists; built from eigenproblems on $A^T A$ (and AA^T). Underlies PCA, low-rank approximation, and regularized regression.

Not required for Math Camp — names and connections only.

Self-study video. <https://www.youtube.com/watch?v=PFDu9oVAE-g>

3.8 Definitions — quadratic form

Setup. Symmetric $A \in \mathbb{R}^{n \times n}$ (Hessians, covariance matrices, $A^T A$).

Four types (via the **quadratic form** $\mathbf{v}^T A \mathbf{v}$):

- **PSD** (positive semidefinite): $\mathbf{v}^T A \mathbf{v} \geq 0$ for all $\mathbf{v}$
- **PD** (positive definite): $\mathbf{v}^T A \mathbf{v} > 0$ for all $\mathbf{v} \neq \mathbf{0}$
- **NSD** (negative semidefinite): $\mathbf{v}^T A \mathbf{v} \leq 0$ for all $\mathbf{v}$
- **ND** (negative definite): $\mathbf{v}^T A \mathbf{v} < 0$ for all $\mathbf{v} \neq \mathbf{0}$

3.8 Examples — where PSD and NSD appear

$A^T A$ is **PSD**. For any $\mathbf{v}$,

$$\mathbf{v}^T A^T A \mathbf{v} = \|A\mathbf{v}\|^2 \geq 0.$$

Covariance matrix Σ is PSD. For centered data, $\Sigma = \mathbb{E}[\mathbf{x}\mathbf{x}^T]$, so

$$\mathbf{v}^T \Sigma \mathbf{v} = \mathbb{E}[(\mathbf{v}^T \mathbf{x})^2] \geq 0$$

(an average of squares).

3.8 Example — the Hessian

Reminder (Day 2). For $f : \mathbb{R}^n \rightarrow \mathbb{R}$ in C^2 , the **Hessian** $H_f(\mathbf{x})$ is the matrix of **second partial derivatives**:

$$(H_f(\mathbf{x}))_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j}(\mathbf{x}).$$

In 2 variables: $H_f = \begin{pmatrix} f_{x_1 x_1} & f_{x_1 x_2} \\ f_{x_2 x_1} & f_{x_2 x_2} \end{pmatrix}$.

Symmetric. Mixed partials agree ($f_{x_i x_j} = f_{x_j x_i}$), so $H_f(\mathbf{x})^\top = H_f(\mathbf{x})$ — a symmetric matrix, so Section 3.8 applies.

Next module (Day 4). We connect H_f to **concavity**: concave $f \Leftrightarrow H_f(\mathbf{x})$ is **NSD** for all $\mathbf{x}$. Today: NSD is the sign condition on the quadratic form $\mathbf{v}^\top H_f(\mathbf{x}) \mathbf{v}$.

3.8 Characterizations — eigenvalues and diagonal entries

Assume $A = A^T$.

Eigenvalues (if $A\mathbf{v}_i = \lambda_i\mathbf{v}_i$) — **necessary** conditions:

- **PSD** $\Rightarrow$ every $\lambda_i \geq 0$
- **PD** $\Rightarrow$ every $\lambda_i > 0$
- **NSD** $\Rightarrow$ every $\lambda_i \leq 0$
- **ND** $\Rightarrow$ every $\lambda_i < 0$

For symmetric A , these are also **sufficient** (equivalent to the definitions); proof omitted at Math Camp.

Diagonal entries ($A = (a_{ij})$) — **necessary** conditions: $\text{PSD} \Rightarrow a_{ii} \geq 0$; $\text{NSD} \Rightarrow a_{ii} \leq 0$. If A is **diagonal**, these sign checks are also **sufficient**.

3.8 Leading principal minors — definition

Leading principal minor. For $A \in \mathbb{R}^{n \times n}$, the k th leading principal minor is

$$\Delta_k = \det(\text{top-left } k \times k \text{ block of } A), \quad k = 1, \dots, n.$$

Example (3×3). $A = \begin{pmatrix} a & b & c \\ d & e & f \\ g & h & i \end{pmatrix}$:

$$\Delta_1 = a, \quad \Delta_2 = \det \begin{pmatrix} a & b \\ d & e \end{pmatrix} = ae - bd, \quad \Delta_3 = \det(A).$$

Only the **north-west** blocks — not e.g. the bottom-right 2×2 block unless it lines up from row 1.

3.8 Leading principal minors — PSD, PD, NSD, ND

Assume $A = A^T$. Leading principal minors give **equivalent** tests in the strict cases; for PSD use **all** principal minors.

- **PD:** $\Delta_k > 0$ for $k = 1, \dots, n$ (Sylvester)
- **PSD:** every **principal minor** ≥ 0 (leading alone is **not** enough)
- **NSD:** $(-1)^k \Delta_k \geq 0$ for $k = 1, \dots, n$
- **ND:** $(-1)^k \Delta_k > 0$ for $k = 1, \dots, n$

Proofs of principal-minor criteria omitted at Math Camp level.

3.8 2×2 checklist

$$A = \begin{pmatrix} a & b \\ b & c \end{pmatrix}, \Delta_1 = a, \Delta_2 = ac - b^2.$$

- **PSD:** $a \geq 0, c \geq 0, ac - b^2 \geq 0$
- **PD:** $a > 0, ac - b^2 > 0$
- **NSD:** $a \leq 0, c \leq 0, ac - b^2 \geq 0$
- **ND:** $a < 0, ac - b^2 > 0$

Example. $A = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$: $a, c > 0, ac - b^2 = 3 > 0 \Rightarrow$ PD (matches eigenvalues 3, 1 from Section 3.7).

3.8 Exercise: 3×3 — leading principal minors Take home

Exercise. For

$$A = \begin{pmatrix} 2 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 2 \end{pmatrix},$$

compute $\Delta_1, \Delta_2, \Delta_3$ and decide whether A is PD, PSD only, or neither.

Solution. $\Delta_1 = 2 > 0$. $\Delta_2 = \det \begin{pmatrix} 2 & -1 \\ -1 & 2 \end{pmatrix} = 4 - 1 = 3 > 0$.


$\Delta_3 = \det(A) = 2(4 - 1) + (-1)(-2) = 6 + 2 = 8 > 0$. All leading principal minors are positive $\Rightarrow A$ is **PD** (Sylvester, Section 3.8).

Appendix

Not required for Math Camp. Proofs and extra examples linked from the main slides.

Order matches main text: A.1–A.2 (Section 3.2), A.3 (3.5), A.4 (3.6), A.5–A.11 (3.7).

Proof idea: closest pair, then a line $\perp \mathbf{p}$



A.1 Proof. Separation when A is compact

for interest only

Claim. Disjoint nonempty convex $A, B \subset \mathbb{R}^n$ with A compact and B closed. Then $\exists \mathbf{p} \neq \mathbf{0}, \beta \in \mathbb{R}$ with $\mathbf{p} \cdot \mathbf{a} \leq \beta \leq \mathbf{p} \cdot \mathbf{b}$ for all $\mathbf{a} \in A, \mathbf{b} \in B$.

Step 1 — closest pair. Define $d = \inf\{\|\mathbf{a} - \mathbf{b}\| : \mathbf{a} \in A, \mathbf{b} \in B\} > 0$. Pick sequences $(\mathbf{a}_n) \subset A, (\mathbf{b}_n) \subset B$ with $\|\mathbf{a}_n - \mathbf{b}_n\| \rightarrow d$. Since A is compact, a subsequence $\mathbf{a}_{n_k} \rightarrow \mathbf{a}_0 \in A$. The $(\mathbf{b}_{n_k})$ are bounded, so (using B closed) a further subsequence $\mathbf{b}_{n_{k_j}} \rightarrow \mathbf{b}_0 \in B$. Hence the infimum is **attained** at some pair $(\mathbf{a}_0, \mathbf{b}_0) \in A \times B$.

Not an interior point. $\mathbf{a}_0$ and $\mathbf{b}_0$ need only lie in A and B ; they are usually on the **boundaries**, not in $\text{int}(A)$ or $\text{int}(B)$.

Step 2 — normal vector. Set $\mathbf{p} = \mathbf{b}_0 - \mathbf{a}_0 \neq \mathbf{0}$ (points from $\mathbf{a}_0$ toward $\mathbf{b}_0$).

A.1 Step 1 — why the subsequence argument works

for interest only

Why $a_{n_k} \rightarrow a_0 \in A$. In $\mathbb{R}^n$, **compact** = closed + bounded (Heine–Borel). Every sequence in a compact set has a **subsequence** converging to a limit **in that set**. So $(a_n) \subset A$ compact $\Rightarrow \exists a_{n_k} \rightarrow a_0 \in A$.

Why (b_{n_k}) is bounded. Since $\|a_{n_k} - b_{n_k}\| \rightarrow d$ and (a_{n_k}) converges, for large k ,

$$\|b_{n_k}\| \leq \|b_{n_k} - a_{n_k}\| + \|a_{n_k}\| \leq d + 1 + \|a_0\| + 1.$$

(Triangle inequality: $\|b\| \leq \|b - a\| + \|a\|$.)

Why $b_{n_{k_j}} \rightarrow b_0 \in B$. A bounded sequence in $\mathbb{R}^n$ has a convergent subsequence (Bolzano–Weierstrass). Since B is **closed**, the limit b_0 lies in B .

A.1 Step 3 — why the derivative works

for interest only

Setup. Fix $a \in A$. For $t \in [0, 1]$, convexity gives $a_t = a_0 + t(a - a_0) \in A$ and $g(t) = \|a_t - b_0\|^2$. Since (a_0, b_0) minimizes distance, $t = 0$ minimizes g on $[0, 1]$.

Why use the derivative? A minimum at the left endpoint $t = 0$ means the slope cannot point downward: $g'(0) \geq 0$.

Differentiate. $g(t) = \|a_0 + t(a - a_0) - b_0\|^2$, so

$$g'(t) = 2\langle a_0 + t(a - a_0) - b_0, a - a_0 \rangle, \quad g'(0) = 2\langle a_0 - b_0, a - a_0 \rangle \geq 0.$$

Hence $\langle b_0 - a_0, a - a_0 \rangle \leq 0$. With $\mathbf{p} = b_0 - a_0$: $\mathbf{p} \cdot (a - a_0) \leq 0$, so $\mathbf{p} \cdot a \leq \mathbf{p} \cdot a_0$.

A.1 Proof. Separation when A is compact (continued)

Step 4 — B stays on the other side. Fix $b \in B$. For $t \in [0, 1]$, set $b_t = b_0 + t(b - b_0) \in B$ and $h(t) = \|a_0 - b_t\|^2$. Since $t = 0$ minimizes h on $[0, 1]$, the same derivative argument gives $\mathbf{p} \cdot b \geq \mathbf{p} \cdot b_0$ for all $b \in B$.

Step 5 — finish. Since $A \cap B = \emptyset$, the closest points satisfy $a_0 \neq b_0$. With $\mathbf{p} = b_0 - a_0$,

$$\mathbf{p} \cdot b_0 - \mathbf{p} \cdot a_0 = \mathbf{p} \cdot (b_0 - a_0) = \|b_0 - a_0\|^2 = d^2 > 0,$$

so $\mathbf{p} \cdot a_0 < \mathbf{p} \cdot b_0$. Any β with $\mathbf{p} \cdot a_0 \leq \beta \leq \mathbf{p} \cdot b_0$ works (e.g. $\beta = \mathbf{p} \cdot a_0$ gives **weak** separation: $\mathbf{p} \cdot a \leq \beta \leq \mathbf{p} \cdot b$ for all $a \in A, b \in B$).

A.1 Strict separation

for interest only

Weak vs. strict. From Steps 3–5: $\mathbf{p} \cdot a \leq \mathbf{p} \cdot a_0$ and $\mathbf{p} \cdot b \geq \mathbf{p} \cdot b_0$. So any β with $\mathbf{p} \cdot a_0 \leq \beta \leq \mathbf{p} \cdot b_0$ separates weakly (e.g. $\beta = \mathbf{p} \cdot a_0$).

When is separation strict? If $d = \|b_0 - a_0\| > 0$ (sets bounded apart; e.g. both compact and disjoint), then $\mathbf{p} \cdot a_0 < \mathbf{p} \cdot b_0$ with a **positive gap**. Pick β **strictly between** them, e.g. $\beta = \frac{1}{2}(\mathbf{p} \cdot a_0 + \mathbf{p} \cdot b_0)$. Then

$$\mathbf{p} \cdot a \leq \mathbf{p} \cdot a_0 < \beta < \mathbf{p} \cdot b_0 \leq \mathbf{p} \cdot b \quad \Rightarrow \quad \mathbf{p} \cdot a < \beta < \mathbf{p} \cdot b.$$

If $d = 0$ (sets touch in the limit, e.g. $A = [0, 1]$, $B = (1, 2]$), only weak separation is possible.

A.1 Proof. Remarks

for interest only

If only B is compact. Swap the roles of A and B , apply the proof, then relabel the normal vector.

Single point $B = \{b_0\}$. Let $\hat{a}$ be a point in A closest to b_0 (exists since A is closed and convex). The same calculation with $(a_0, b_0) = (\hat{a}, b_0)$ separates A from b_0 even when A is not compact.

A.2 Proof. Efficient production $\Rightarrow$ profit maximization

for interest only

Goal. $y \in Y$ technically efficient $\Rightarrow \exists \mathbf{p} \neq \mathbf{0}$ with $\mathbf{p} \cdot y \geq \mathbf{p} \cdot y'$ for all $y' \in Y$.

Assumptions (camp level). $Y \subset \mathbb{R}^n$ is **closed and convex**; Y is **bounded** (so compact); **free disposal** (if $y \in Y$ and $\tilde{y} \leq y$, then $\tilde{y} \in Y$).

Step 1 — better set. Define

$$P_y = \{y' \in \mathbb{R}^n : y' \geq y, y' \neq y\}$$

(weakly more output, strictly more in at least one good). P_y is **convex**.

Step 2 — disjointness. If $y' \in P_y \cap Y$, then y' weakly dominates y in Y , contradicting efficiency. So $P_y \cap Y = \emptyset$.

A.2 Proof. Efficient production (continued)

for interest only

Step 3 — separate. Y is compact and P_y is closed, so apply the separating hyperplane theorem (swap labels if needed so a compact set is on one side). There exist $\mathbf{p} \neq \mathbf{0}$ and β with

$$\mathbf{p} \cdot \mathbf{y}'' \leq \beta \leq \mathbf{p} \cdot \mathbf{y}' \quad \forall \mathbf{y}'' \in Y, \forall \mathbf{y}' \in P_y.$$

Hence $\mathbf{p} \cdot \mathbf{y}'' \leq \mathbf{p} \cdot \mathbf{y}'$ for all $\mathbf{y}'' \in Y, \mathbf{y}' \in P_y$.

Step 4 — profit maximization at y . For each good i , pick $\mathbf{y}^{(i)} = y + \varepsilon \mathbf{e}_i \in P_y$ for small $\varepsilon > 0$ (feasible with free disposal). Then $\mathbf{p} \cdot \mathbf{y}'' \leq \mathbf{p} \cdot \mathbf{y}^{(i)} = \mathbf{p} \cdot y + \varepsilon p_i$. Let $\varepsilon \downarrow 0$ to get $\mathbf{p} \cdot \mathbf{y}'' \leq \mathbf{p} \cdot y$ for all $\mathbf{y}'' \in Y$.

Remark. Separation gives only $\mathbf{p} \neq \mathbf{0}$. For $\mathbf{p} \geq 0$, need free disposal on all goods: if $p_i < 0$, then $y - \varepsilon \mathbf{e}_i \in Y$ (since $\leq y$) raises profit, contradicting maximization at y .

A.3 Transpose rules — proof

for interest only

$(A^T)^T = A$. Transpose swaps rows and columns twice, so you return to A .

$(AB)^T = B^T A^T$. Entry (i, j) of $(AB)^T$ equals entry (j, i) of AB :

$$((AB)^T)_{ij} = (AB)_{ji} = \sum_k a_{jk} b_{ki} = \sum_k (B^T)_{ik} (A^T)_{kj} = (B^T A^T)_{ij}.$$

Same matrix $\Rightarrow$ equal.

$(A\mathbf{x})^T = \mathbf{x}^T A^T$. Special case with $B = \mathbf{x}$ a column $(n \times 1)$: $(A\mathbf{x})^T = \mathbf{x}^T A^T$.

A.3 Inverse rules — proof

for interest only

$(AB)^{-1} = B^{-1}A^{-1}$ (both invertible). Multiply in reverse order:

$$(AB)(B^{-1}A^{-1}) = A(BB^{-1})A^{-1} = AIA^{-1} = I.$$

Similarly $(B^{-1}A^{-1})(AB) = I$. So $B^{-1}A^{-1}$ is the inverse of AB .

$(A^{-1})^{-1} = A$. By definition of inverse, $A^{-1}A = I$, so A is the inverse of A^{-1} .

$(A^{-1})^T = (A^T)^{-1}$. Transpose $A^{-1}A = I$:

$$(A^{-1}A)^T = A^T(A^{-1})^T = I.$$

So $(A^{-1})^T$ is the inverse of A^T .

A.4 Projection matrix — why $P_X^2 = P_X$?

for interest only

Setup. $P_X = X(X^T X)^{-1} X^T$ is the **projection matrix** (hat matrix); $M_X = I - P_X$ is the **annihilator matrix** (residual maker matrix).

Idempotent: $P_X^2 = P_X$ (**project once is enough**). $P_X \mathbf{y} = \hat{\mathbf{y}}$ already lies in $\text{Col}(X)$. Vectors in that yellow plane are unchanged by projecting onto it again:

$$P_X \hat{\mathbf{y}} = \hat{\mathbf{y}} \quad \implies \quad P_X(P_X \mathbf{y}) = P_X \mathbf{y}.$$

First drop the part orthogonal to $\text{Col}(X)$; a second drop removes nothing.

A.4 Dot product and transpose — why $(A\mathbf{u}) \cdot \mathbf{v} = \mathbf{u} \cdot (A^T\mathbf{v})$?

for interest only

Goal. This is the **transpose rule** used in the symmetry proof of P_X .

Rewrite both sides as row $\times$ column. Dot product = transpose the first vector, then multiply:

$$(A\mathbf{u}) \cdot \mathbf{v} = (A\mathbf{u})^T \mathbf{v} = \mathbf{u}^T A^T \mathbf{v},$$

using $(A\mathbf{u})^T = \mathbf{u}^T A^T$.

Other side is the same expression.

$$\mathbf{u} \cdot (A^T \mathbf{v}) = \mathbf{u}^T (A^T \mathbf{v}).$$

So $(A\mathbf{u}) \cdot \mathbf{v} = \mathbf{u} \cdot (A^T \mathbf{v})$ for every $A, \mathbf{u}, \mathbf{v}$.

Meaning: A^T moves A from the left of $\mathbf{u}$ to the right of $\mathbf{u}$ inside an inner product.

A.4 Projection matrix — why $P_X^T = P_X$?

for interest only

Step 1 — geometric balance. $P_X \mathbf{u} \in \text{Col}(X)$, so it is orthogonal to the part of $\mathbf{v}$ outside that space:

$$(P_X \mathbf{u}) \cdot \mathbf{v} = (P_X \mathbf{u}) \cdot (P_X \mathbf{v}).$$

Swapping $\mathbf{u} \leftrightarrow \mathbf{v}$ gives the same with $\mathbf{u}$ on the right. Hence

$$(P_X \mathbf{u}) \cdot \mathbf{v} = \mathbf{u} \cdot (P_X \mathbf{v}) \quad \text{for all } \mathbf{u}, \mathbf{v} \in \mathbb{R}^n.$$

Step 2 — transpose rule. Previous slide: $(P_X \mathbf{u}) \cdot \mathbf{v} = \mathbf{u} \cdot (P_X^T \mathbf{v})$.

Step 3 — conclude symmetry. Steps 1 and 2 give $\mathbf{u} \cdot (P_X^T \mathbf{v}) = \mathbf{u} \cdot (P_X \mathbf{v})$ for all $\mathbf{u}, \mathbf{v}$. Fix $\mathbf{v} = \mathbf{e}_j$ and $\mathbf{u} = \mathbf{e}_i$: the j th column of P_X^T equals the j th column of P_X for every j . So $P_X^T = P_X$.

Direct check: $(P_X)^T = X((X^T X)^{-1})^T X^T = P_X$, using $(X^T X)^{-1} = ((X^T X)^T)^{-1} = (X^T X)^{-1}$.

A.5 Symmetric matrices — eigenvalues are real

for interest only

Claim. If $A = A^T$ (real symmetric), every eigenvalue λ is **real**.

Proof (setup). Let $A\mathbf{v} = \lambda\mathbf{v}$ with $\mathbf{v} \neq \mathbf{0}$. We allow **complex** $\mathbf{v}$: over $\mathbb{C}$ eigenvalues always exist; we show $\lambda \in \mathbb{R}$.

$\mathbf{v}^H =$ **conjugate transpose** of $\mathbf{v}$ (transpose, then conjugate each entry).

Multiply $A\mathbf{v} = \lambda\mathbf{v}$ on the left by $\mathbf{v}^H$:

$$\mathbf{v}^H A \mathbf{v} = \lambda \mathbf{v}^H \mathbf{v}.$$

Here $\mathbf{v}^H \mathbf{v} = \|\mathbf{v}\|^2 > 0$ (a positive real number).

A.5 Symmetric matrices — eigenvalues are real (cont.)

for interest only

Left side is real. Conjugate-transpose rule: $(\mathbf{v}^H \mathbf{A} \mathbf{v})^* = \mathbf{v}^H \mathbf{A}^H \mathbf{v}$. A is real and symmetric, so $A^H = A^T = A$:

$$(\mathbf{v}^H \mathbf{A} \mathbf{v})^* = \mathbf{v}^H \mathbf{A} \mathbf{v} \implies \mathbf{v}^H \mathbf{A} \mathbf{v} \in \mathbb{R}.$$

(A scalar equals its own conjugate $\Leftrightarrow$ it is real.)

Finish. From $\mathbf{v}^H \mathbf{A} \mathbf{v} = \lambda \mathbf{v}^H \mathbf{v}$, divide by $\mathbf{v}^H \mathbf{v} > 0$:

$$\lambda = \frac{\mathbf{v}^H \mathbf{A} \mathbf{v}}{\mathbf{v}^H \mathbf{v}} \in \mathbb{R}.$$

Real numerator, positive real denominator $\Rightarrow \lambda \in \mathbb{R}$.

Not all symmetric eigenvalues are positive — e.g. $A = \text{diag}(1, -1)$. Sign conditions (PSD/NSD) are in Section 3.8.

A.6 Equivalent linear systems

for interest only

Equivalent systems. Two linear systems (same unknowns) are **equivalent** if they have **exactly the same solutions**.

Easier equivalent system. Same solution set, but written in a form that is simpler to read — e.g. fewer redundant equations, leading variables isolated, zeros below each pivot. Row reduction replaces a system by an easier equivalent one **without changing who solves it**.

Homogeneous case. For $M\mathbf{x} = \mathbf{0}$, every equivalent system has the same **kernel** (null space)

$$\ker M = \{\mathbf{x} : M\mathbf{x} = \mathbf{0}\}.$$

Finding eigenvectors for λ is the same problem with $M = A - \lambda I$.

A.6 Elementary row operations

for interest only

Allowed operations on a matrix (or augmented matrix $[M \mid \mathbf{b}]$):

1. **Swap** two rows.
2. **Scale** one row by a **nonzero** number.
3. **Add** a multiple of one row to another row.

Each operation produces an **equivalent system** (same solution set).

Why solutions are preserved (idea). Swapping or scaling an equation does not change its truth set. If row i and row j are both satisfied, so is $(\text{row } j) + c (\text{row } i)$; conversely you can undo the combination using the other allowed ops.

RREF. Keep applying these ops until the matrix is in **reduced row echelon form**: leading 1 in each pivot row; only zero above/below each leading 1; each leading 1 lies to the right of the one above.

A.6 Pivot variables and free variables

for interest only

Once a matrix is in **RREF**, each column is either a **pivot column** or a **free column**.

Pivot column / pivot variable. A column with a **leading** 1; the matching unknown is a **pivot variable**. Its RREF row expresses the pivot variable as an **affine function of the free variables**.

Free column / free variable. A column *without* a leading 1; the matching unknown is a **free variable**. You may set each free variable to **any** value; nothing in RREF forces it.

Key step. Read each pivot row and **write each pivot variable as a function of the free variables**:

$$x_{\text{pivot}} = (\text{expression involving only free variables}).$$

Once the free variables are chosen, every pivot variable is determined; the whole solution is fixed.

A.6 Example: row reduction (non-trivial)

for interest only

Problem. Find all $\mathbf{x} \in \mathbb{R}^4$ with $M\mathbf{x} = \mathbf{0}$ for

$$M = \begin{pmatrix} 1 & 0 & 2 & 1 \\ 2 & 0 & 4 & 2 \\ 0 & 1 & 1 & 0 \\ 0 & 2 & 2 & 0 \end{pmatrix}.$$

(Coupled equations; not block-diagonal.)

Row ops. $R_2 \leftarrow R_2 - 2R_1$ eliminates the duplicate x_1 -row; $R_4 \leftarrow R_4 - 2R_3$ eliminates the duplicate x_2 -row:

$$\longrightarrow \begin{pmatrix} 1 & 0 & 2 & 1 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix} \xrightarrow{\text{swap } R_2, R_3} \begin{pmatrix} 1 & 0 & 2 & 1 \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix} \quad (\text{RREF}).$$

A.6 Example: pivot variables as functions of free variables

for interest only

From the RREF above: pivot columns 1, 2 $\Rightarrow$ pivot variables x_1, x_2 ; free columns 3, 4 $\Rightarrow$ free variables x_3, x_4 .

Equivalent system:

$$x_1 + 2x_3 + x_4 = 0, \quad x_2 + x_3 = 0.$$

Solve each pivot row for its pivot variable:

$$x_1 = -2x_3 - x_4, \quad x_2 = -x_3, \quad x_3, x_4 \in \mathbb{R} \text{ free.}$$

Every solution is fixed once (x_3, x_4) is chosen.

A.6 Basis of the kernel from free variables

for interest only

General solution. Substitute the free variables into $\mathbf{x} = (x_1, x_2, x_3, x_4)^\top$:

$$\mathbf{x} = x_3 \begin{pmatrix} -2 \\ -1 \\ 1 \\ 0 \end{pmatrix} + x_4 \begin{pmatrix} -1 \\ 0 \\ 0 \\ 1 \end{pmatrix}.$$

So $\ker M = \text{span}\{\mathbf{b}_1, \mathbf{b}_2\}$ with $\dim \ker M = 2$ (two free variables).

Recipe (basis of $\ker M$). For each free variable in turn:

1. set that free variable = 1;
2. set all other free variables = 0;
3. compute pivot variables from the RREF rows.

Here: $(x_3, x_4) = (1, 0) \Rightarrow \mathbf{b}_1 = (-2, -1, 1, 0)^\top$; $(x_3, x_4) = (0, 1) \Rightarrow \mathbf{b}_2 = (-1, 0, 0, 1)^\top$.

Check: $M\mathbf{b}_1 = \mathbf{0}$, $M\mathbf{b}_2 = \mathbf{0}$.

A.7 When eigenvalues are not real

for interest only

The characteristic equation $\det(A - \lambda I) = 0$ may have **no real roots**. Then A has **no real eigenvectors** (no real direction kept by scaling alone).

Example: 90° **rotation** (from earlier slides), $A = \begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix}$:

$$\det(A - \lambda I) = \lambda^2 + 1 = 0 \quad \Rightarrow \quad \lambda = \pm i \text{ (not real).}$$

Every nonzero vector is rotated off its own line, so no $\mathbf{v} \in \mathbb{R}^2 \setminus \{\mathbf{0}\}$ satisfies $A\mathbf{v} = \lambda\mathbf{v}$.

Over $\mathbb{C}$. Allow complex λ and $\mathbf{v}$; the same recipe (characteristic polynomial, then $(A - \lambda I)\mathbf{v} = \mathbf{0}$) works. For **real** A , non-real eigenvalues come in **conjugate pairs** $\lambda, \bar{\lambda}$.

Existence over $\mathbb{C}$: **appendix (for interest only)**. When a real eigenbasis fails, see A.9–A.11 for QR and SVD (**for interest only**).

A.8 Orthogonal and orthonormal matrices

for interest only

Orthonormal columns. A square matrix $Q \in \mathbb{R}^{n \times n}$ is **orthogonal** if

$$Q^T Q = I \quad (\text{equivalently } Q Q^T = I).$$

Column i is $\mathbf{q}_i$; then $\mathbf{q}_i^T \mathbf{q}_j = \delta_{ij}$ (orthonormal).

Geometric meaning. Q is a **rigid motion** of $\mathbb{R}^n$ (rotation and/or reflection):

$$\|Q\mathbf{x}\| = \|\mathbf{x}\|, \quad (Q\mathbf{x})^T (Q\mathbf{y}) = \mathbf{x}^T \mathbf{y}.$$

Lengths and angles are preserved; $Q^{-1} = Q^T$.

Terminology. “Orthonormal matrix” usually means the same as orthogonal (unit-length orthonormal columns).

A.8 QR and SVD (brief)

for interest only

Same spirit as eigen-decomposition: rewrite A so the action is easy to read.

- **QR** ($A = QR$): Q has orthonormal columns, R is upper triangular. Used for stable least squares and picking independent columns ([appendix](#)).
- **SVD** ($A = U\Sigma V^T$): always exists; U, V orthogonal, Σ diagonal (singular values). Scales along orthogonal directions — works even when real eigenvectors do not exist (e.g. rotation).

More detail on main slide “Decompositions — QR, eigen, SVD” (for interest only). Existence: [appendix](#), [appendix](#) (**for interest only**).

A.9 Existence of eigenvalues and eigenvectors (over $\mathbb{C}$)

for interest only

Claim. Every $A \in \mathbb{C}^{n \times n}$ has at least one eigenpair $(\lambda, \mathbf{v})$ with $\mathbf{v} \neq \mathbf{0}$.

Step 1 (characteristic polynomial). $p_A(\lambda) = \det(A - \lambda I)$ is a polynomial in λ of degree n .

Step 2 (eigenvalue exists). By the **fundamental theorem of algebra**, $p_A(\lambda) = 0$ has n roots in $\mathbb{C}$ (counting multiplicity). Pick one root λ .

Step 3 (eigenvector exists). $p_A(\lambda) = 0$ means $\det(A - \lambda I) = 0$, so $A - \lambda I$ is **singular**. Hence $(A - \lambda I)\mathbf{v} = \mathbf{0}$ has a nontrivial solution $\mathbf{v} \neq \mathbf{0}$.

Real A : non-real roots come in conjugate pairs $\lambda, \bar{\lambda}$.

A.10 Existence of QR — what the decomposition means

for interest only

Claim. Every $A \in \mathbb{R}^{m \times n}$ can be written as $A = QR$: Q has **orthonormal columns**, R is **upper triangular**.

Setup. Write $A = [\mathbf{a}_1 \ \mathbf{a}_2 \ \mathbf{a}_3 \ \cdots]$ — a **series of column vectors** 1, 2, 3, ... They may overlap (not perpendicular, different lengths).

Goal (Gram–Schmidt). Build new **clean directions** $\mathbf{q}_1, \mathbf{q}_2, \mathbf{q}_3, \dots$ (perpendicular, unit length) so that
vector j is built only from $\mathbf{q}_1, \dots, \mathbf{q}_j$.

The numbers r_{ij} in R record **how much** of each vector goes along each direction.

A.10 Existence of QR — vectors 1 and 2

for interest only

Vector 1. Take the first column, normalize it — that is the first clean direction:

$$\mathbf{q}_1 = \frac{\mathbf{a}_1}{\|\mathbf{a}_1\|}.$$

Vector 1 lies entirely along $\mathbf{q}_1$.

Vector 2. Project $\mathbf{a}_2$ onto $\mathbf{q}_1$, subtract, and keep the **residual** as the new direction:

$$\mathbf{a}_2 = \underbrace{(\mathbf{q}_1^T \mathbf{a}_2) \mathbf{q}_1}_{\text{part along } \mathbf{q}_1} + \underbrace{(\mathbf{a}_2 - (\mathbf{q}_1^T \mathbf{a}_2) \mathbf{q}_1)}_{\perp \mathbf{q}_1}.$$

Normalize the residual to get $\mathbf{q}_2$. **Key:** vector 2 uses **only** $\mathbf{q}_1$ and $\mathbf{q}_2$: $\mathbf{a}_2 = r_{12}\mathbf{q}_1 + r_{22}\mathbf{q}_2$.

A.10 Existence of QR — vectors 3, 4, ...

for interest only

Vector 3. Project $\mathbf{a}_3$ onto $\mathbf{q}_1$ and $\mathbf{q}_2$, subtract both overlaps. Normalize what's left $\Rightarrow \mathbf{q}_3$. Vector 3 uses **only** $\mathbf{q}_1, \mathbf{q}_2, \mathbf{q}_3$.

Vector 4, ... Same pattern: for vector k , subtract its projections onto all previous clean directions, normalize the residual $\Rightarrow \mathbf{q}_k$.

General step.

$$\mathbf{a}_k = \sum_{j=1}^{k-1} (\mathbf{q}_j^T \mathbf{a}_k) \mathbf{q}_j + (\text{residual } \perp \text{ all previous } \mathbf{q}_j).$$

So $\mathbf{a}_k = r_{1k}\mathbf{q}_1 + \cdots + r_{kk}\mathbf{q}_k$ — no $\mathbf{q}_\ell$ with $\ell > k$ needed.

A.10 Existence of QR — reading R and $A = QR$

for interest only

What is r_{ij} ? From $\mathbf{a}_j = \sum_i \mathbf{q}_i r_{ij}$: entry r_{ij} is **how much of vector j lies along direction i :**

$$r_{ij} = \mathbf{q}_i^T \mathbf{a}_j.$$

Why upper triangular? Vector j is built from $\mathbf{q}_1, \dots, \mathbf{q}_j$ only, so $r_{ij} = 0$ for $i > j$. Diagonal $r_{jj} = \text{length of the residual}$ (the genuinely new part).

Existence. Run the process on all columns; stack $\mathbf{q}_j$ into Q and r_{ij} into R . Then $A = QR$.

A.11 Existence of SVD — intuition

for interest only

Dimensions. $A \in \mathbb{R}^{m \times n}$ maps $\mathbb{R}^n \rightarrow \mathbb{R}^m$: $\mathbf{x} \in \mathbb{R}^n \Rightarrow A\mathbf{x} \in \mathbb{R}^m$.

Why SVD? Eigenvectors need $A\mathbf{v} = \lambda\mathbf{v}$ (same direction in and out) — fine for square A with a full eigenbasis. Fails for rotation (no real eigenvectors) and when $m \neq n$ (input and output live in different spaces).

Idea. Find **unit** input directions $\mathbf{v}_i$ ($\|\mathbf{v}_i\| = 1$) and unit output directions $\mathbf{u}_i$ ($\|\mathbf{u}_i\| = 1$) with

$$A\mathbf{v}_i = \sigma_i\mathbf{u}_i, \quad \sigma_i \geq 0.$$

Input direction $\mathbf{v}_i$ maps to output direction $\mathbf{u}_i$, stretched by σ_i (no mixing).

Picture. Unit sphere in $\mathbb{R}^n \rightarrow$ ellipsoid in $\mathbb{R}^m$. $\mathbf{v}_i$ = input axes; $\mathbf{u}_i$ = where they land; σ_i = axis lengths. Factorization: $A = U\Sigma V^T$ = rotate input (V), scale (Σ), rotate output (U).

A.11 Existence of SVD — derivation (Step 1)

for interest only

Step 1 — input directions ($\mathbb{R}^n$). Diagonalize $A^T A$ ($n \times n$): orthonormal eigenvectors $\mathbf{v}_i$ with $\|\mathbf{v}_i\| = 1$, eigenvalues $\lambda_i \geq 0$. Set $\sigma_i = \sqrt{\lambda_i} = \|\mathbf{A}\mathbf{v}_i\|$ (stretch along unit input $\mathbf{v}_i$).

Remark. λ_i is real (A.5, [appendix](#)) and $\lambda_i \geq 0$ (Section 3.8) because $A^T A$ is symmetric PSD. So $\sigma_i = \sqrt{\lambda_i}$ is real and ≥ 0 ; it is **not** an eigenvalue of A .

A.11 Existence of SVD — derivation (Steps 2–3)

for interest only

Step 2 — output directions ($\mathbb{R}^m$). Goal: $A\mathbf{v}_i = \sigma_i \mathbf{u}_i$ with $\|\mathbf{u}_i\| = 1$. From Step 1: $\mathbf{v}_i$ is **normalized** ($\|\mathbf{v}_i\| = 1$) and $\sigma_i = \|A\mathbf{v}_i\|$.

Apply A to the unit input $\mathbf{v}_i$: the vector $A\mathbf{v}_i \in \mathbb{R}^m$ points the right way but has length σ_i . Normalize it:

$$\mathbf{u}_i = \frac{1}{\sigma_i} A\mathbf{v}_i \quad \implies \quad A\mathbf{v}_i = \sigma_i \mathbf{u}_i, \quad \|\mathbf{u}_i\| = 1.$$

Step 3 — assemble. $V = [\mathbf{v}_1 \cdots]$, $U = [\mathbf{u}_1 \cdots]$, $\Sigma = \text{diag}(\sigma_1, \dots) \Rightarrow A = U\Sigma V^\top$.